

Indian English and Telugu speech dataset for TTS

This is a 60-minute speech dataset for training text-to-speech models: 30.2 minutes of Indian English and 30.1 minutes of Telugu, cut from YouTube recordings, transcribed and emotion-tagged with Sarvam's APIs, and published on HuggingFace. This report describes the pipeline, the data-quality changes that came out of reviewing the output, and the checks run on the result.

- Dataset: <https://huggingface.co/datasets/AkCodes23/sarvam-tts-in-te-en>
- Code: <https://github.com/AkCodes23/Sarvam-AI>

Every clip contains a single voice, tracked by `speaker_id`. I deliberately kept several speakers per language (4 English, 5 Telugu, 9 in total) to add accent and speaking-style variety rather than build a one-voice corpus. Most of the sources are storytelling, so the dataset is largely Indian narrative speech: mythology, folk tales, and audiobook fiction.

In line with the brief's emphasis on listening to the data, a manual listening audit of 80 clips (40 English, 40 Telugu), sampled across all speakers and source types, checked transcript accuracy, speaker consistency, and audio cleanliness. The results are in section 8.

1. What I built and how the pipeline works

The pipeline is a modular Python package (`ttsds`) driven by a CLI. It processes one source at a time, so each source's output could be checked and API spend kept down before scaling up. The stages:

1. **Source curation.** I hand-picked YouTube sources that are single-voice by nature: solo audiobooks, storytelling channels, single-narrator lectures, and one stage talk. Compilation channels and anything with a music bed were left out at this step. Source selection had the largest effect on downstream quality, so it got the most manual review.
2. **Download and normalize.** `yt-dlp` pulls the best-quality audio plus provenance (channel, date, license, video id). `ffmpeg` produces a 16 kHz copy for ASR and a 24 kHz mono master for the published audio.
3. **Batch ASR with diarization** (Sarvam `saaras:v3`). This returns speaker-labelled, time-stamped chunks, which is what makes it possible to keep only the stretches from one speaker.
4. **Segmentation.** Neighbouring chunks from the target speaker are merged into longer runs, and a run ends wherever a different speaker appears. Each run is then split into clips of 3 to 25 seconds, with a target of 5 to 15. Sarvam's chunk boundaries are only approximate, so each cut is snapped to the nearest silence using local energy, which keeps clips from starting or ending in the middle of a word.
5. **Per-segment ASR** (Sarvam `saarika:v2.5`). The batch transcript is tied to the coarse diarization chunks, so once a clip is cut it is transcribed again on its own. This second pass gives a transcript that lines up with the actual clip, along with word timings for trimming and a per-clip language-confidence score. It doubles the ASR cost, but it is the main reason the per-clip transcripts hold up at the word level.
6. **Acoustic features** (`parselmouth`, `librosa`). For each clip the pipeline measures pitch, energy and how much it varies, HNR, spectral tilt, voicing fraction, speaking rate, and pauses. These are z-scored within each speaker, so what counts as "excited" is judged against how that speaker normally sounds.
7. **Quality gates.** Per clip, with an explicit reason recorded: sustained clipping, low SNR, excessive silence, music or noise bed (inter-pause energy ratio), low ASR confidence, implausible character rate, near-duplicate transcript, wrong detected language, and a code-mix flag for a Telugu clip that is majority English. A clip that fails a gate is dropped; softer concerns are kept as flags for review.
8. **Emotion and style tagging.** A mix of acoustic features and an LLM (Sarvam `sarvam-30b`, the smaller of the two models used here). The per-speaker acoustic profile is written out as text and sent with the


transcript under a fixed taxonomy. The whispered-speech style is set by an acoustic rule rather than letting the LLM infer it from the words. Each tag carries a confidence and a source (auto or human).

9. **Topic and validation.** A second call per clip, to the larger sarvam-105b, assigns a topic and independently judges the transcript and the emotion label (section 4).
10. **Human review.** A static HTML app lists every candidate with its audio, transcript, tag, and metrics, with buttons to accept, reject, relabel, or fix a transcript. A human edit takes priority over the automated label.
11. **Balance, finalize, publish.** Selection prefers storytelling topics, then the cleanest clips by DNSMOS, and balances the emotion histogram to about 30 minutes per language. Audio gets light loudness normalization with no hard limiting, which preserves the loudness swings that carry emotion. The dataset is built with HuggingFace datasets (two configs, each with train/validation/test) and pushed public.

Final dataset:

Metric	Indian English	Telugu
Minutes	30.2	30.1
Clips (train / val / test)	160 (144 / 8 / 8)	150 (134 / 8 / 8)
Speakers	4	5
Storytelling clips	79%	75%
Median DNSMOS	3.08	3.16
Clips above DNSMOS 3.0	57%	81%
Median alignment confidence	0.95	0.94

The emotion mix is balanced on purpose. The five common labels (neutral, calm, sad, excited, angry) are each capped near 27 to 30 clips so none of them dominates, and the rarer ones are kept in full wherever they appear. Telugu ends up with all eight labels; English has seven, since no English clip read as surprised (it has 6 happy and 8 fearful, but 0 surprised). Beyond emotion, every row carries a normalized transcript, language, style, gender and accent, the quality scores (DNSMOS, SQUIM, SNR, alignment, intra-clip cohesion), a topic, the validator's suitability score, full source provenance, and timestamps.



Schema. Each row in the published dataset has these fields:

Column	Description
audio	24 kHz mono waveform
text	raw Sarvam transcript
normalized_text	TTS-facing text (numbers and abbreviations expanded)
language, language_code	en / te, en-IN / te-IN
emotion, style	emotion label (8-class) and speaking style (narrative, conversational, formal, expressive)
emotion_confidence, tag_source	tagger confidence 0–1; auto or human
speaker_id, gender, accent	speaker identifier and inferred attributes
duration	clip length in seconds
dnsmos_ovrl (+ sig, bak), dnsmos_pass	DNSMOS perceptual quality and the >3.0 gate
snr_db, squim_stoi/pesq/sisdr	SNR and reference-free quality estimates
mms_align_score, overlap_flag	forced-alignment confidence; possible second voice
topic, llm_tts_suitable	topic category; validator suitability 0–1
annotation_flags (+ booleans)	quality_flag, has_truncation, has_codemix, transcript_review_needed, ...
source_video_id, source_url, source_channel, license	provenance
segment_start, segment_end, sample_rate	source timing and sample rate

A real row (en_mahabharata_0025):

Field	Value
text	"If you ever break your promise, I'll leave you at once."
emotion / style	sad / narrative
emotion_confidence	0.85
speaker_id / gender / accent	en_mahabharata / female / Indian English
duration	4.29 s
dnsmos_ovrl / snr_db / mms_align_score	3.49 / 46.5 dB / 0.97
topic	fiction

2. Iterations to improve data quality

Each change below came out of reviewing the pipeline's intermediate output.

- Clipping gate rejected clean audio.** The first Telugu audiobook lost 41 of 43 clips to a "clipping" gate even though the audio was fine. The clips peaked at exactly 1.0 because YouTube masters loud, and the gate had treated any sample at full scale as clipping. Actual clipping shows up as runs of flat-topped samples, so the gate was changed to measure the fraction of flat-topped samples instead, which recovered all 43 clips.
- The LLM copied the prompt template.** Every clip came back "neutral, narrative" with the rationale field reading " ≤ 20 words", which was the instruction text. The model was completing the template literally instead of reasoning about the clip. Removing the template and describing the fields instead spread the labels out, and they then matched the audio.
- The reasoning model truncated before answering.** Labels were varied but half had low confidence. sarvam-30b reasons before answering, and at a 1500-token budget it spent all of it reasoning and never wrote the JSON. Raising the budget (billing is by tokens actually used, so a higher ceiling costs nothing) stopped the truncation, and median tag confidence then reached 0.85.
- DNSMOS re-curation.** Perceptual quality (DNSMOS) flagged 47% of clips below 3.0, over the pre-set one-third threshold. Per-source analysis showed three sources dragging the set down: an archival AIR recording (2.31), a hall discourse (2.27), and an English audiobook (2.42) whose compression DNSMOS heard even though its SNR looked clean. Dropping all three and adding a clean lecture and more narration raised the pool pass rate from 53 to 63 percent.
- Topic focus.** The sources were mostly storytelling, so that became the dataset's theme rather than a random mix, with LLM topic tags used to prefer narrative clips during selection.
- Telugu transcript check switched to an Indic recognizer.** The first independent recognizer for Telugu was generic Whisper, but it disagreed with Sarvam at 76 percent word error, which reflected Whisper's weakness on Telugu rather than transcript errors. A Telugu-specialized recognizer brought that down to 47 percent, and forced alignment and the listening pass became the primary Telugu checks (sections 3 and 4).
- Validation moved to a larger model.** The first emotion cross-check effectively had the tagging model grade its own taxonomy. Validation moved to the larger sarvam-105b judging the smaller sarvam-30b's tags, which is a more independent test than self-grading (sections 4 and 7).
- The low-confidence emotion flag was over-firing.** An early version of emotion_low_confidence folded in the judge's endorsement and flagged about 60 percent of clips, which made it useless. Keying

it to the tag's own confidence (below 0.55) instead left only the genuinely uncertain clips flagged, 0 in English and 2 in Telugu.

9. **An annotation layer for imperfect-but-usable clips.** Reading and listening to clips surfaced recurring imperfections: truncated utterances, mild background noise, the occasional possible second voice. Rather than drop those clips, an annotation layer (`quality_flag`, `has_truncation`, `transcript_review_needed`, `overlap_flag`) was added so they ship labeled and filterable (section 5).

3. What worked and what did not

Worked:

- Audiobook and storytelling sources gave the best mix of clean audio and emotional range.
- The double-pass ASR produced clip-accurate transcripts. English cross-checks at 5.8% word error against an independent recognizer.
- DNSMOS exposed bad sources that SNR alone missed (the compressed audiobook), and per-source numbers made it possible to drop those specific sources instead of cutting across the board.
- ECAPA speaker verification confirmed single-speaker integrity (AUC 0.96).
- MMS forced alignment provides a transcript check that holds up in Telugu. Switching the second recognizer to an Indic, Telugu-specialized model roughly halved the Telugu cross-ASR word error versus generic Whisper (76 percent down to 47 percent).

Did not work or needed care:

- Emotion labels are weakly supervised. Off-the-shelf SER models cluster toward neutral and do not transfer to Telugu, and the larger model endorses 26 percent of the smaller model's labels on acoustic grounds. The labels ship with confidences and are best used as an expressive-TTS filter that a human can refine, not as ground truth.
- pyannote overlap detection is behind a license that cannot be accepted from a script, so an intra-clip embedding-cohesion check stands in for it.
- Telugu cross-ASR stays inconclusive even with the Indic recognizer: it still disagrees with Sarvam at 47 percent word error, since two independent Telugu systems diverge heavily on spelling and word splits. For Telugu, forced alignment and the listening pass carry more weight than cross-ASR.
- Clean studio-grade Indian English is scarce on YouTube. The cleanest English clips are off-topic (a lecture, a talk), which forced a trade-off described in section 4.


4. Quality observations and decisions

Each observation below is backed by a measurement, and where it helps, a specific clip is cited that can be pulled up in the dataset to check it.

Evaluation methodology. Every number in this report comes from a script in `scripts/`, run over the published clips, with the output saved as JSON under `data/manifests/`. Speaker verification embeds each clip with an ECAPA-TDNN model (SpeechBrain) and scores 10,000 balanced pairs (5,000 same-speaker, 5,000 different-speaker) by cosine similarity; the ROC-AUC and equal-error rate come from those scores (`eval_speaker_eer.py`). Perceptual quality is Microsoft DNSMOS run per clip, with SQUIM for STOI/PESQ/SI-SDR (`score_audio_quality.py`). Alignment confidence is the mean per-frame probability from the torchaudio MMS forced aligner (`score_mms_align.py`). Cross-ASR word error is computed with `jiwer` against an independent recognizer (`eval_asr.py`, `eval_asr_indic.py`). Emotion reliability compares the 30b tags against an independent 105b re-labelling on a 120-clip stride sample (60 per language), reporting percent agreement, Cohen's κ , and Krippendorff α (`eval_emotion.py`,


`eval_agreement.py`). Phoneme coverage uses grapheme-to-phoneme conversion (`g2p_en` for English ARPAbet, `epitran` for Telugu IPA) over the normalized transcripts, and lexical diversity (type-token ratio, Zipf) is computed over the same text (`eval_phoneme.py`, `phoneme_freq.py`).


Single-speaker. Every clip is embedded with ECAPA-TDNN and the embeddings compared by cosine similarity. Same-speaker pairs score 0.74 on average and different-speaker pairs 0.21. Treated as a verification task over 10,000 pairs, that separation gives an AUC of 0.96 and a 9 percent equal-error rate, so clips from one speaker do not bleed into another's.




Transcripts. For English, an independent recognizer (`Whisper small.en`) agrees with the Sarvam transcripts at 5.8 percent word error, and the realtime recognizer identified the correct language on 100 percent of clips. For Telugu, an Indic recognizer was used rather than generic Whisper: a Telugu-specialized Whisper fine-tune (`vasista22`) brings the word error against Sarvam to 47 percent (median 50), down from 76 percent with generic `Whisper-large`. That is still high: two independent Telugu ASR systems disagree heavily on a morphologically rich, sandhi-heavy script. So for Telugu I weight MMS forced alignment (median 0.94 confidence, English 0.95) and the human listening pass above cross-ASR. Word error is reported throughout rather than character error, since it is the stricter test of whether the words are right.


The second ASR pass is what catches the worst transcript errors. On `en_mahabharata_e67_0019`, for example, the coarse batch pass produced "Karna about to creeper", and the per-clip pass corrected it to "turned about Kripa". Most clips need no correction: `en_mahabharata_0025` transcribes cleanly as "If you ever break your promise, I'll leave you at once." The full set of edge cases and their resolutions is in section 5.






Emotion. The emotion labels are best read as weakly supervised annotations, intended for filtering expressive TTS data, rather than as ground-truth emotion. Their reliability was measured rather than assumed. On a 120-clip sample (60 per language), the smaller tagging model (sarvam-30b) and the larger validation model (sarvam-105b) gave the same label 65 percent of the time, with Cohen's κ 0.55 and Krippendorff α 0.44 between them, which is moderate agreement. The disagreement is structured, not random (heatmap below): `calm` is the least stable label, scattering to `sad` (8 clips) and `neutral` (6), and every `happy` clip was relabelled `excited`, while `sad` and `excited` are firmly on the diagonal. Two acoustic speech-emotion models (emotion2vec, audeering) drop the three-rater Krippendorff α to about 0.01, because they collapse toward neutral and were not trained for Telugu. In a separate pass where the larger model judges whether the acoustics support the 30b label, it endorses 26 percent of them. Because the labels are weakly supervised, each one ships with a confidence and a source field. Low-confidence clips are flagged (`emotion_low_confidence`), and a user can keep only the high-confidence or only the expressive clips by filtering on those fields. For instance `en_mahabharata_0040` ("The angry king said, I don't care for my promise") is tagged angry at 0.95 confidence, while `te_ramaaraavi_0040` is tagged angry at only 0.40 and would be removed by a confidence threshold. Laughter, where it occurs, is recorded as an emotion (usually happy or excited) rather than as a separate flag.






Overlap. pyannote is gated, so short windows within each clip were embedded and checked to confirm they all match the same speaker. Median cohesion is 0.58, the single-speaker level, and only 9 clips fall below 0.40 and are flagged for a listen.

Perceptual quality and the topic trade-off. DNSMOS scores how clean audio sounds. Selecting for the storytelling topic, the published set is 57 percent above 3.0 in English and 81 percent in Telugu. English is lower because its cleanest clips are a lecture and a talk, which are off-topic, so preferring storytelling pulls in narration at DNSMOS 2.85 to 3.19. I chose topic consistency over marginal gains in perceptual quality, since the set is more useful as a focused storytelling corpus, and the `dnsmos_pass` column still recovers the cleanest subset with one filter. These clips sit at 23 to 35 dB SNR and the validator rated them around 0.90 suitable, so the lower DNSMOS comes from how the audio was recorded and mastered, not from added noise. One example is `en_mahabharata_e67_0035` ("Right. No other human being can defeat Karna."), at DNSMOS 2.94, 23 dB SNR, and 0.90 suitability.



Per-source quality. Emotion entropy measures how varied a source's emotions are (higher is more varied).

Source	Type	Clips	Min	Median SNR	Median DNSMOS	Emotion entropy
en_nptel	lecture	53	6.9	35.7	3.23	2.36
en_mahabharata	story	50	8.5	35.4	3.19	2.60
en_air_talk	talk	44	9.9	28.6	3.02	2.24
en_tedx_amina	talk	46	9.2	43.9	2.97	2.03
en_mahabharata_e69	story	40	8.1	24.0	2.88	2.52
en_mahabharata_e67	story	40	8.1	22.9	2.85	2.39
te_ramaaraavi	story	47	8.6	44.9	3.23	2.34
te_kalalavelugu	audiobook	50	9.4	30.6	3.20	2.60
te_kathalu_epic	story	39	8.8	39.9	3.16	2.52
te_motivation_kasyap	talk	47	7.9	38.6	3.01	2.31
te_bhumiputri	audiobook	43	10.0	27.5	2.95	2.53



Selection funnel. After the per-clip gates, the kept pool was 499 clips (273 English, 226 Telugu). The published release is 310 of those (160 English, 150 Telugu), chosen to reach about 30 minutes per language while keeping the emotion histogram balanced. The other 189 passed every check but were balanced out: once a language hit its 30-minute target, surplus clips from the over-represented buckets (mostly neutral and calm) were dropped rather than let one emotion dominate. Few clips were rejected at the gate stage itself, because the bad recordings had already been removed earlier, at source curation and the DNSMOS step; the gates are the last line of defence, not the first. Concrete error and edge-case examples are in section 5.

Other decisions. Light loudness normalization was used instead of a hard target, so the intensity dynamics that carry emotion survive. Human edits always override automated labels. Gender is inferred from median F0 with known speakers corrected by hand. The full 60 minutes is kept rather than applying a hard DNSMOS cut, with `dnsmos_pass` exposing the studio-grade subset.

5. Edge cases and annotation decisions

Real speech is messier than a scripted studio read, so most clips carry some imperfection. Every clip carries an `annotation_flags` string and matching boolean fields that record what is imperfect, and users filter on them.

Flag definitions. - `quality_flag`: an automatically inferred audio-quality concern, not a verified audible-noise label (set when DNSMOS OVRL is under 3.0, or SNR under 18 dB, or there is elevated energy in the pauses). I named it `quality_flag`, not `has_noise`, because it does not assert that the clip contains audible noise. - `low_quality_audio`: the stricter automated quality proxy, DNSMOS OVRL under 2.8 (clearly degraded). - `has_truncation`: the transcript does not end on terminal punctuation, so the clip likely ends mid-utterance. - `has_codemix`: the regional-language clip contains preserved English words (see the convention below). - `emotion_low_confidence`: the emotion tag's own confidence is below 0.55. - `transcript_review_needed`: the LLM judge flagged the transcript, or MMS alignment is below 0.85. - `overlap_flag`: intra-clip speaker cohesion is low, a possible second voice.

Statistics (published set).

Flag	English (of 160)	Telugu (of 150)
<code>quality_flag</code>	68	28
<code>has_truncation</code>	11	31
<code>transcript_review_needed</code>	17	46
<code>low_quality_audio</code>	36	10
<code>overlap_flag</code>	3	4
<code>emotion_low_confidence</code>	0	2
<code>has_codemix</code>	0	0

Conventions.

Laughter is treated as an emotion rather than its own flag. When a clip is audibly laughing, it gets labelled by the emotion behind it, usually happy or excited. The narration sources here do not contain laughter, so in practice this never comes up. Other non-verbal sounds (noise, cough, breath, a long pause) can be marked inline with tags like `<cough>` when they are clearly audible, but those come from a listening pass and are not in the auto-generated transcripts yet.

When a clip is cut off mid-sentence, it is marked in `annotated_text` with a trailing em dash, as in "...does it have to do with —", while the raw `text` field is left clean for training.

Code-mixing would normally be handled by keeping the English word in Latin script and bracketing it, for example "aa [project] inka [complete] kaaledu". In this dataset it never triggers. Sarvam's ASR transliterates English into Telugu script, so an English word like "project" comes back written in Telugu characters, and the detector that looks for Latin-script runs finds none. `has_codemix` is therefore zero across all 150 Telugu clips. The handling is in place, but there is no code-mixing left in the text to annotate; catching genuine code-switches would need an ASR pass built for code-mixed speech, which is left for future work.

Edge-case audit (real examples).


Clip	Issue	Original	Final	Resolution
en_mahabharata_e67_0019	ASR garbled a name	"Karna about to creeper"	"turned about Kripa"	per-clip realtime re-ASR fixed it
AIR talk (ax7l6qOqgpQ)	clip ends mid-utterance	"...does it have to do with"	"...does it have to do with —"	has_truncation=true, em dash in annotated_text
en_air_talk_0024	emotion ambiguous	30b: sad	105b: neutral	advisory; emotion confidence plus flag, kept
en_mahabharata_0004	emotion ambiguous (nearby)	30b: happy	105b: excited	advisory; both plausible, kept
en_mahabharata_e67 clips	DNSMOS 2.85 to 2.9, mild	(kept)	(kept)	quality_flag=true, recoverable via dnsmos_pass filter

Curation decisions. I kept clips with minor issues and annotated them explicitly rather than removing them. Noisy English storytelling clips (DNSMOS just under 3.0) stayed because they carry the storytelling voice and emotion the dataset is built around, and `quality_flag` / `dnsmos_pass` let a user exclude them. Three whole sources that DNSMOS exposed as perceptually poor (2.3 to 2.4) were dropped, because no per-clip flag rescues a bad recording. I let topic coherence outweigh DNSMOS on the English side, accepting a lower clean-rate for a coherent storytelling corpus; a user who disagrees can use the flags to rebuild a cleaner or different subset without touching the audio. The flags also cluster usefully.


`transcript_review_needed` is higher in Telugu (46 vs 17), consistent with Telugu being the harder ASR target, and the AIR talk source accounts for most truncated and emotion-ambiguous English clips, which is why it contributes few clips to the final set.

6. TTS readiness analysis


Phoneme coverage. Coverage is estimated by grapheme-to-phoneme conversion over the normalized transcripts (g2p_en into ARPAbet for English, epitran into IPA for Telugu), counting the unique phonemes that appear against each language's inventory. English covers all 39 ARPAbet phonemes (100 percent); Telugu covers 44 of roughly 50 (about 88 percent). The thinnest Telugu phonemes are the aspirated and breathy-voiced consonants: the retroflex aspirate and breathy-g appear once each, the palatal aspirate twice, and aspirated dental, k, and p between 10 and 13 times each. These are marginal phones, entering Telugu mainly through Sanskrit and loanwords. They are under-represented in any natural corpus this size, so a model trained here will see them rarely. In English the rarest are ZH (0.03 percent) and OY (0.09 percent), as expected.




Duration. Clips run 3.1 to 24.5 seconds in English (median 11.6) and 3.1 to 22.8 in Telugu (median 13.1). Durations spread across the range, with no pile-up at the 3-second floor or the 25-second ceiling.




Transcript length. Median 27 words per clip in English and 26 in Telugu, a comfortable single utterance length.



Speech rate. English centers on 144 words per minute (6 percent slow, 65 percent medium, 29 percent fast); Telugu on 122 (18 percent slow, 80 percent medium, 2 percent fast). A spread of speech rates across the set helps a model generalize across tempos.



Lexical diversity. English has 1,193 unique words over 30 minutes (type-token ratio 0.27), Telugu 1,947 (0.52, higher because Telugu inflects heavily). The word-frequency curve follows the expected Zipf shape, which indicates natural running text with the usual head of common function words, not a small set of repeated lines.



Text normalization. The `normalized_text` field is the TTS-facing reading of each transcript: language-aware expansion of numbers, ordinals, currency, and a conservative set of abbreviations into spoken words, so a model trains on "sixteenth day" and "doctor" instead of "16th" and "Dr.". Numbers run through `num2words` in the clip's own language, so "15" becomes "fifteen" in Telugu, with a safe fallback to the digits if a value is unsupported. Most of the corpus is narrative, so this only edits the handful of clips that carry numerals. The abbreviation and currency rules are there for robustness rather than because the content needs them often. The raw `text` field is kept verbatim alongside it for reference.

How this compares to a typical TTS corpus. Here is how the set lines up against typical TTS training data, without singling out any one corpus.

Property	Common practice	This dataset
Sample rate	22.05 to 24 kHz mono	24 kHz mono
Total size	tens of hours per voice for production work, an hour or two for a prototype	about 60 minutes total
Speakers	often one voice with many hours, or many voices with little each	9 voices, roughly 4 to 7 minutes each
Clip length	usually 1 to 15 seconds	3 to 25 seconds, median about 12
Transcripts	typically hand-typed or hand-corrected	Sarvam ASR, cross-checked and spot-audited by ear
Speaking style	often neutral read speech	expressive narration with emotion tags
Loudness	sometimes normalized hard	light normalization, dynamics kept
Provenance and splits	varies	per-clip source and license, train/validation/test splits

On the things a modern TTS pipeline expects, the set is in line: the sample rate, file format, the train/validation/test splits, and the per-clip provenance all match common practice, and the emotion coverage is wider than the neutral read speech many corpora ship with. Where it differs is size and shape. An hour of audio is prototype scale, not production scale, and spreading it over nine voices leaves only a few minutes per speaker, so this is better suited to multi-speaker or fine-tuning experiments than to training a single high-fidelity voice. The transcripts are ASR-derived rather than hand-typed, which is the reason so much of this report is spent checking them.


7. Dataset analysis

Beyond the per-clip checks in section 4, six dataset-level analyses look at how the data is distributed and where its limits are (`scripts/analyze_dataset.py`, `eval_emotion.py`).

Split integrity. The train/validation/test split is stratified by emotion and speaker. No clip appears in more than one split, no transcript is duplicated across splits, and every speaker in validation and test is also present in training, so there is no leakage and the held-out sets follow the training distribution. The emotion histogram is preserved in each split. Split sizes are 144/8/8 for English and 134/8/8 for Telugu.


Per-speaker distribution and quality. The nine voices are far from evenly sized. On the English side one narrator, en_mahabharata, supplies 23.9 of the 30.2 English minutes (about 79 percent), while the other three English voices contribute 1.6 to 2.5 minutes each. Telugu is more balanced, spread across five speakers of 1.9 to 7.5 minutes. The practical consequence is that the English half behaves close to a single-speaker set, which is worth knowing before training a multi-speaker model.

Speaker	Lang	Gender	Clips	Minutes	Median DNSMOS
en_mahabharata	en	F	122	23.9	2.95
en_nptel	en	M	20	2.5	3.34
en_air_talk	en	F	10	2.2	3.17
en_tedx_amina	en	F	8	1.6	3.29
te_ramaaraavi	te	F	41	7.5	3.24
te_audiobook_kalalavelugu	te	F	41	7.5	3.22
te_kathalu_epic	te	M	30	6.8	3.16
te_audiobook_bhumiputri	te	F	27	6.3	2.97
te_motivation_kasyap	te	M	11	1.9	3.04




Emotion agreement and confusion. The two LLM raters agree on 65 percent of the 120-clip sample (Cohen's κ 0.55), and the disagreement is concentrated on adjacent classes (calm to sad and neutral, happy to excited). The confusion matrix and the full discussion are in section 4.


Valence-arousal consistency. Even where the categorical labels disagree, the acoustic valence-arousal estimates line up with them in the expected direction. Positive labels (happy, surprised, calm) sit at higher valence than negative ones (sad, angry, fearful), and high-arousal labels (angry, excited, fearful) sit above the low-arousal ones (calm, sad, neutral). The labels are therefore at least consistent with the prosody, which is what justifies using them as a weakly supervised filter.



Audio bandwidth and spectral quality. This one is a genuine limitation. Although every clip is stored at 24 kHz, the audio is effectively band-limited: 99 percent of the energy sits below about 4 kHz in English and 2.6 kHz in Telugu, and only 0.1 to 0.2 percent of the energy is above 8 kHz. Some of that is speech's natural low-frequency dominance, but the near-absence of high-frequency content points to YouTube's lossy encoding low-passing the sources. The set is well suited to standard 16 to 24 kHz TTS, but not to full-band, high-fidelity synthesis where crisp sibilance matters.



SNR versus perceptual quality. SNR and DNSMOS are only moderately correlated (Pearson $r = 0.52$), which is why both are kept. A clip can have a clean noise floor (high SNR) yet a low DNSMOS because of codec or mastering artifacts the SNR estimate cannot see; the compressed English audiobook from section 2 was exactly that case.



8. Human quality audit

Quality was checked two ways: automated validation, and a manual listening audit of the clips by ear.

The automatic layer is the larger model (sarvam-105b) reading each clip's transcript and acoustic summary, with no access to the audio. Across the 310 published clips it judged 89 percent of transcripts clean and 81 percent suitable to train on, and supported the smaller model's emotion label on 26 percent. This checks the labels; it does not replace listening to the audio.

The human layer is the listening audit, which the automated raters cannot do. I built the harness (scripts/human_audit.py sample writes a scoring sheet and an audit.html player; score tallies it) and ran both passes by hand: 40 English and 40 Telugu clips, drawn by stratified random sampling across all speakers and source types, listening to each and comparing the voice against the transcript.

Metric	English (n = 40)	Telugu (n = 40)
Exact transcript match	37 (92.5%)	35 (87.5%)
Minor error (pronunciation differences)	3 (7.5%)	5 (12.5%)
Major error	0	0
Audio perceived clean	31 (77.5%)	32 (80.0%)
Minor background noise	9 (22.5%)	8 (20.0%)
Judged unsuitable for TTS	0	0

Transcripts held up well in both languages: 37 of 40 English and 35 of 40 Telugu matched the audio exactly, the rest differed only in minor pronunciation, and there were no major transcription errors on either side. On audio, minor background noise showed up in 9 English and 8 Telugu clips, but the speaker stayed clearly intelligible in every case and not one clip was judged unusable for TTS. The human noise rate is about a fifth of clips in each language. In English that sits well below the automatic `quality_flag` rate (~43 percent); in Telugu it is comparable (~19 percent). That fits how the flag is meant to work: it is a conservative, over-inclusive proxy a consumer can filter on, and a flag firing does not mean the clip is bad.

9. What I would improve given more time

- A human emotion-labeling pass. The transcript-and-audio listening audit is now done for both languages (section 8); emotion is still checked only by the automatic raters, so a human relabeling pass would turn that last proxy into ground truth. The review tool supports it.
- A speech-emotion model that handles Telugu, so the third emotion rater is fair.
- Word-level forced-alignment trimming to tighten clip edges further.
- Background-music separation to rescue otherwise-good clips that carry a light bed.
- A cleaner Indian-English storytelling source, the one ingredient that stayed scarce, so the English set could reach both topic coherence and high perceptual quality.